

## 1. Iteration Time Analysis:

### A.) Profile time taken for 10 iterations :

```
[cs3401.13@rce problem2]$ cat profile_metrics.csv
iteration,total_time,map_time,shuffle_time,reduce_time,io_time
1,.699071653,.200142809,.009216478,.171740394,.021829070
2,1.093970894,.190360870,.009464200,.171110557,.022118983
3,1.103579650,.190807635,.009361694,.170378042,.022201892
4,1.101387144,.204288877,.009166210,.170009458,.022300698
5,1.106438217,.190528156,.009248777,.181801138,.021752108
6,1.090062649,.189172085,.009240302,.172048684,.022044166
7,.965525758,.190290005,.009393922,.159471015,.021947386
8,1.113359892,.201523542,.009110608,.170377814,.021737397
9,1.117379305,.203165720,.009208903,.170339300,.021674876
10,1.106991690,.189777854,.009084010,.183828016,.021778133
```

On the shared test case containing 1000 nodes, these are the profile times for the 10 iterations.

Average values (approximately):

- Total iteration time = **1.07 sec**
- Map time = **0.195 sec (18%)**
- Shuffle time = **0.009 sec (0.84%)**
- Reduce time = **0.172 sec (16%)**
- IO time = **0.022 sec (2%)**

### B.) Overhead per iteration :

Now, Map + Shuffle + Reduce = **0.376 sec**

=> Overhead =  $1.07 - 0.376 = 0.694 \text{ sec}$

=> **65% of time is framework overhead.**

### C.) Time for I/O : **0.022 sec approx**

## 2. Scalability Analysis:

We analyze scaling with:

- $|V|$  (nodes)
- $|E|$  (edges)

### A.) Theoretical Complexity

Each iteration:

#### Map Phase:

For each node:

- emit adjacency list =>  $O(|V|)$
- emit one contribution per outgoing edge =>  $O(|E|)$

So:

$$\text{Map complexity} = O(V + E)$$

#### Shuffle + Sort:

Sort on intermediate data:

$$\text{Intermediate size} \approx O(V + E)$$

$$\text{Sorting complexity: } O((V+E) \log(V+E))$$

This is the dominant cost.

#### Reduce:

$$\text{One reduce per node: } O(V + E)$$

#### Overall Per Iteration:

$$O((V+E) \log(V+E))$$

Total:

$$O(I (V+E) \log(V+E))$$

### B.) Validation :

For 1000 nodes, with average degree of 5 :

```

iteration,total_time,map_time,shuffle_time,reduce_time,io_time
1,.688740260,.202591640,.008507463,.171395372,.020382741
2,1.101295390,.188100353,.008737501,.170695784,.021165567
3,1.077677635,.188216075,.008594518,.171568174,.020948184
4,1.110663099,.200595707,.008847840,.181602799,.021574595
5,1.099566657,.188788630,.008724266,.172272852,.021507376
6,1.087750103,.200932390,.008890018,.170618654,.021392625
7,1.117459323,.201171549,.009114308,.170955311,.021599475
8,1.093666239,.189473568,.008794205,.172403472,.021263592
9,1.103729576,.188454798,.008712652,.170588673,.021076925
10,1.109676198,.199504188,.008930582,.182233105,.021393708

```

For 5000 nodes, with average degree of 5 :

```

iteration,total_time,map_time,shuffle_time,reduce_time,io_time
1,.764524009,.227314698,.027739465,.191205850,.053911685
2,1.169696276,.217224450,.028478715,.190365537,.054745309
3,1.174209831,.231204841,.028893959,.165428836,.055579227
4,1.188148406,.231383409,.029101086,.191728096,.056074182
5,1.171361402,.229056765,.028930411,.165941346,.055475953
6,1.198880681,.230986456,.029274296,.189240915,.056487459
7,1.197189838,.218973481,.029331380,.189933848,.056544421
8,1.173918462,.219147739,.029193774,.178035636,.056452980
9,1.181059952,.229563133,.029350930,.177887033,.056416459
10,1.171091921,.229834127,.029173217,.166798828,.056133805

```

For 1000 nodes, with average degree of 5 :

```

iteration,total_time,map_time,shuffle_time,reduce_time,io_time
1,.843810128,.258593768,.052345043,.187158462,.094314926
2,1.254217691,.261099525,.052965135,.186568546,.097090864
3,1.250770931,.242995369,.055371262,.187252483,.100542760
4,1.262319208,.253769091,.054683296,.186205052,.099619054
5,1.274638793,.267745201,.055344687,.186525796,.100595531
6,1.250058071,.254448698,.054815529,.186108002,.099862502
7,1.279074868,.266273789,.054950604,.198304371,.099555814
8,1.251817054,.242003903,.055140104,.187460984,.100205700
9,1.257227240,.267139489,.055148441,.199163939,.099988343
10,1.260479813,.266429353,.055234602,.174910938,.099978620

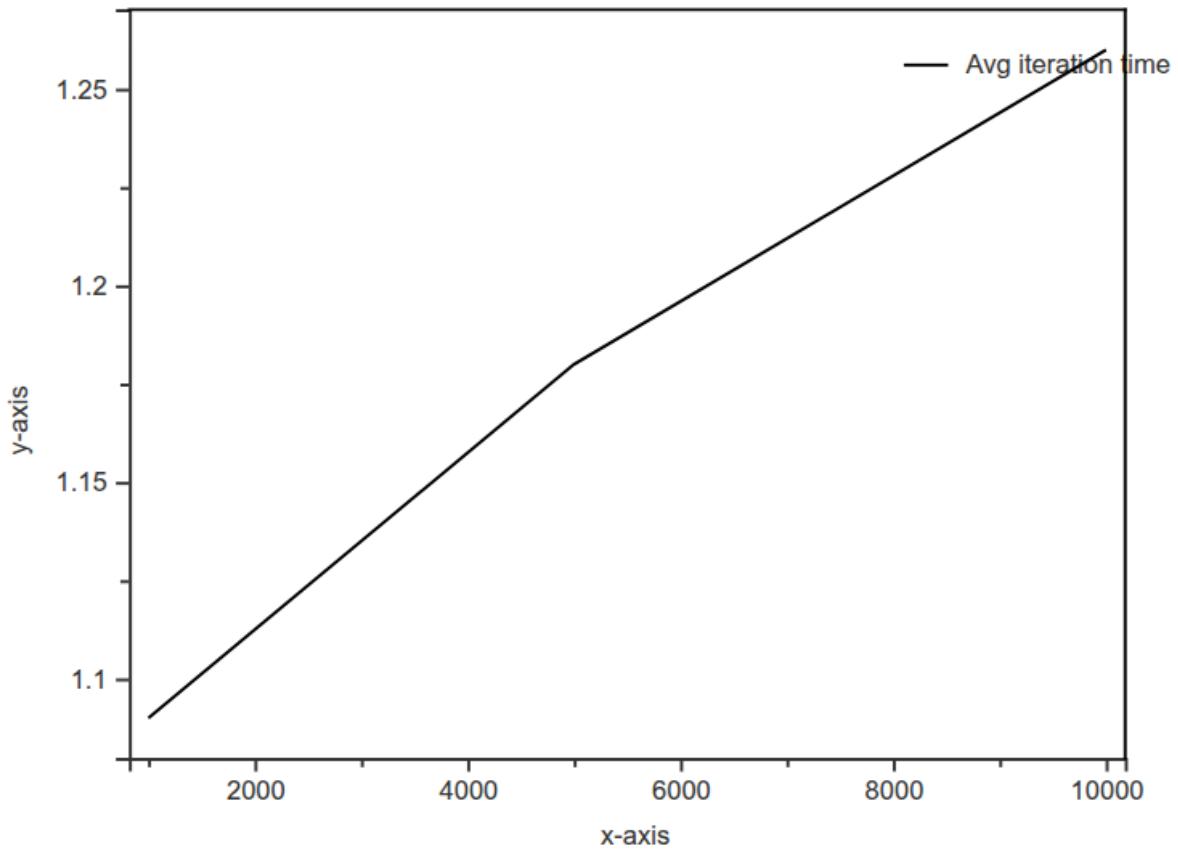
```

=> Average iteration time

for 1000 nodes is **1.09 seconds**

for 5000 nodes is **1.18 seconds**

for 10000 nodes is **1.26 seconds**



### C.) Observations:

From 1k to 10k:

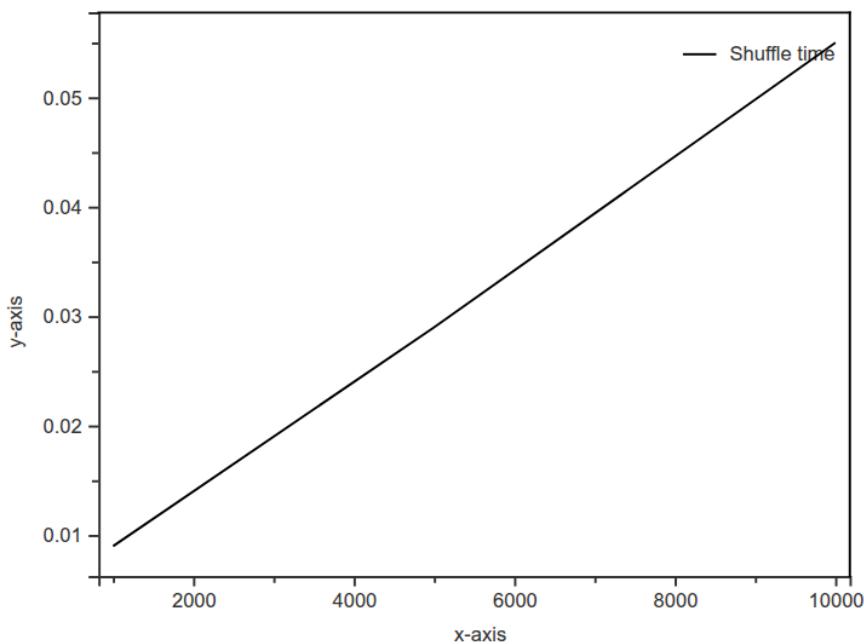
- Nodes increased **10x**
- Edges increased **10x**
- Time increased only **~15%**

That means: runtime is not dominated by graph size yet. It is dominated by fixed overhead.

### Shuffle time :

Nodes	Shuffle
1k	~0.009s
5k	~0.029s
10k	~0.055s

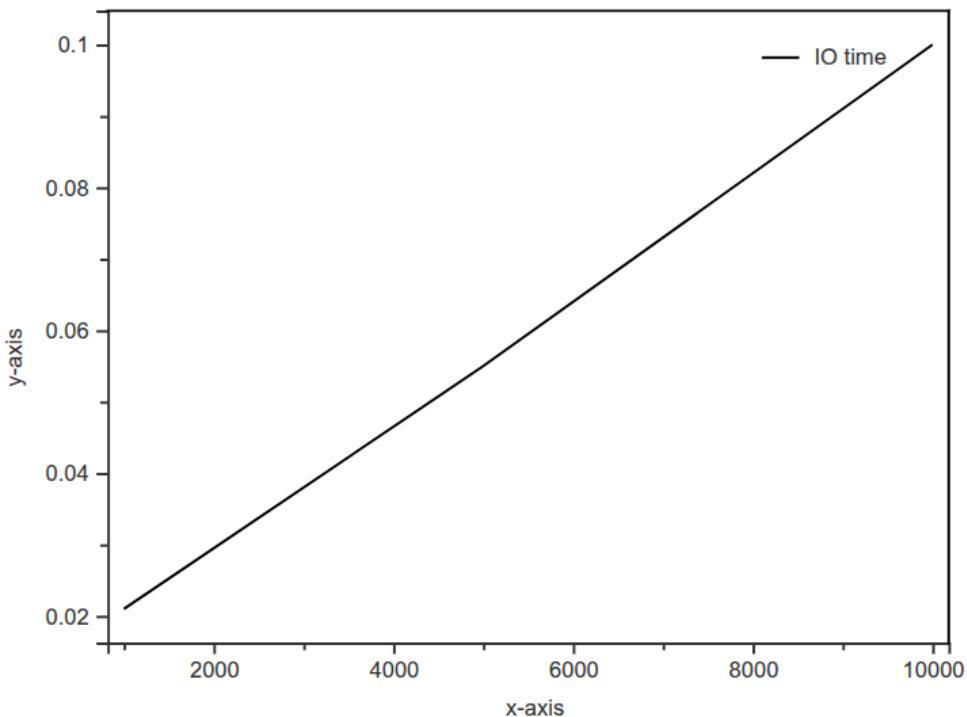
**This scales almost perfectly linearly with edges.** That is expected because shuffle size =  $O(E)$



### IO Time:

Nodes	IO
1k	0.021s
5k	0.055s
10k	0.10s

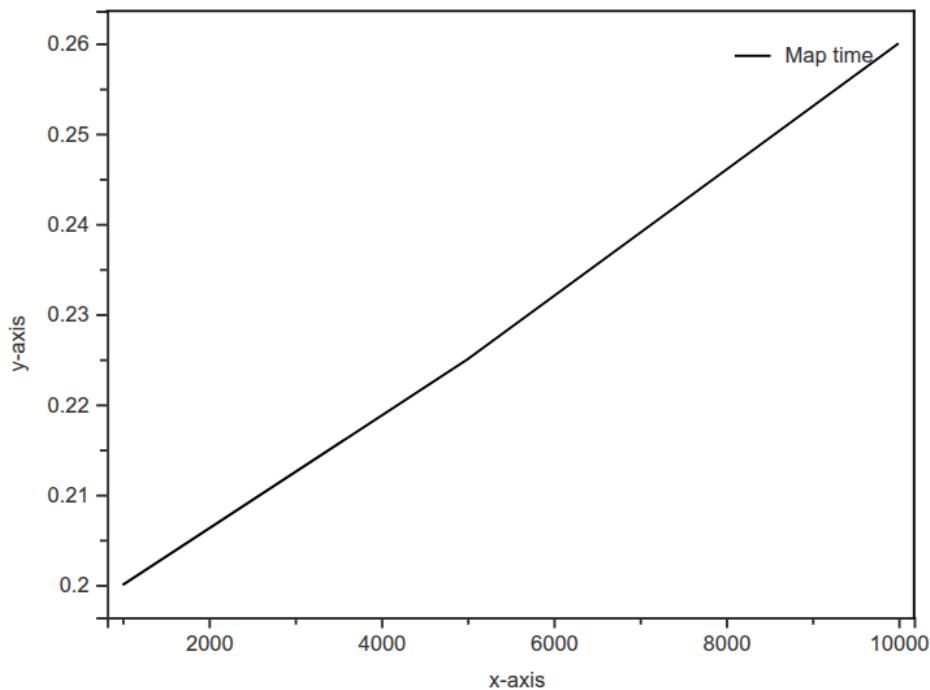
Also nearly linear with E.



### Map Time:

Nodes	Map
1k	0.19–0.20
5k	0.22–0.23

Nodes	Map
10k	0.25–0.27



### Reduce Time:

Almost constant (~0.18–0.19). That means reducer cost is small relative to overhead.

### D.) Compare With Theoretical Complexity

Theoretical per iteration:  $O((V+E)\log(V+E))$

Since avg degree is constant (5):  $E=O(V)$

So:  $T(V) \approx O(V\log V)$

Now check growth: If pure  $V \log V$  -

For 1k:  $1000\log(1000)$

For 10k:  $10000\log(10000)$

Ratio  $\approx 100000\log100000/1000\log1000$

$\approx 13.3\times$

But the observed ratio:

$1.26/1.09 \approx 1.15\times$

This huge difference might be due to the dominance of overhead at small scale, i.e – the true compute time is being hidden by SLURM startup, Barrier sync, Process creation, etc.

**Validation to this claim :** Extract Real Compute Scaling -

Subtract fixed overhead.

From 1k:

Total  $\approx 1.09$

Map + Shuffle + Reduce + IO  $\approx 0.20 + 0.009 + 0.17 + 0.02 \approx 0.40$

So overhead  $\approx 0.69$  sec.

Subtract that overhead from others:

For 5k:

$1.18 - 0.69 \approx 0.49$  sec compute

For 10k:

$1.26 - 0.69 \approx 0.57$  sec compute

Now compare compute only:

$0.40$  to  $0.57 \approx 1.4\times$

Much closer to theoretical scaling.

## E.) Estimating Maximum Graph Size for 10 Minutes

10 iterations of 10k:

$10 \times 1.26 \approx 12.6$  sec

Time limit = 600 sec

Scale factor:

$600 / 12.6 \approx 47\times$

So estimated maximum:

Nodes  $\approx 10k \times 47 \approx \text{470k nodes}$

Edges  $\approx 50k \times 47 \approx \text{2.35 million edges}$

Before hitting 10 min limit.

\*But, sorting will eventually dominate, so actual max slightly lower.

## F.) Conclusions

### Observation 1

Execution time increases sub-linearly with graph size for small graphs because distributed framework overhead dominates computation.

### Observation 2

Shuffle and I/O times scale linearly with number of edges, confirming O(E) behavior.

### **Observation 3**

Sorting cost grows but remains small relative to SLURM startup overhead for graphs <= 10k nodes.

### **Observation 4**

True asymptotic behavior will approach  $O(V \log V)$  once graph size becomes large enough for compute to dominate overhead.

## **3. I/O Performance**

```
iteration,total_time,map_time,shuffle_time,reduce_time,io_time,data_written_bytes
1,.857885915,.259474532,.051899915,.185965274,.094195851,3801431
2,1.258600327,.263630842,.052998031,.186289373,.096936120,4079377
3,1.263931891,.253699119,.054985757,.187341370,.099586610,4456782
4,1.245088458,.255197069,.054247927,.197665426,.099512785,4413192
5,1.265259412,.255682555,.055603319,.199030260,.101227070,4424312
6,1.287056976,.255926970,.055633717,.197778187,.100360434,4434942
7,1.266993689,.255162341,.056403029,.186366156,.101572283,4422517
8,1.228603684,.255073014,.054922233,.175278143,.100828233,4422412
9,1.261404791,.254992419,.055884247,.198851325,.100782001,4426207
10,1.150590617,.253871461,.055553086,.187626833,.100811295,4423792
```

A.) Average data volume per iteration = **4.4MB**

(this is for graph input with 10k nodes, 5 edges on average)

B.) Disk throughput =  $\text{data\_written}/\text{io\_time} = 4.4 \text{ MB}/0.1 \text{ s} = \text{44 MB/s}$

C.) Network overhead - In real distributed Hadoop, shuffle causes network transfer of intermediate key-value pairs. Our approximation can be - Network cost  $\approx$  size of merged.txt. Because in real cluster, that would be transferred across nodes.