

1. Parallel Efficiency Analysis:

Tested on the 1000 nodes, 3000 edges sample test case given.

For 1 process :

```
[cs3401.13@rce problem3_2]$ sbatch run_2node_mpi.sbatch large_file.txt 1
Submitted batch job 56431
[cs3401.13@rce problem3_2]$ ls
compile.sh  large_file.txt  mpi_56431.err  mpi_56431.out  mpi_graph  mpi_gr.
[cs3401.13@rce problem3_2]$ cat mpi_56431.out
=====
SLURM_JOB_ID = 56431
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 0 running on node06.local
Total Execution Time: 8.0799e-05 seconds
0 0
```

For 2 processes :

```
[cs3401.13@rce problem3_2]$ sbatch run_2node_mpi.sbatch large_file.txt 2
Submitted batch job 56432
[cs3401.13@rce problem3_2]$ ls
compile.sh  large_file.txt  mpi_56431.err  mpi_56431.out  mpi_56432.err  mp
[cs3401.13@rce problem3_2]$ cat mpi_56432.out
=====
SLURM_JOB_ID = 56432
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 0 running on node06.local
Process 1 running on node06.local
Total Execution Time: 0.000188627 seconds
0 0
1 0
2 0
3 0
4 0
```

For 4 processes :

```
[cs3401.13@rce problem3_2]$ sbatch run_2node_mpi.sbatch large_file.txt 4
Submitted batch job 56433
[cs3401.13@rce problem3_2]$ ls
compile.sh  large_file.txt  mpi_56431.err  mpi_56431.out  mpi_56432.err  mp
[cs3401.13@rce problem3_2]$ cat mpi_56433.out
=====
SLURM_JOB_ID = 56433
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 1 running on node06.local
Process 0 running on node06.local
Process 3 running on node07.local
Process 2 running on node07.local
Total Execution Time: 0.0018343 seconds
0 0
1 0
2 0
```

Here, we observe that the execution time is increasing instead of decreasing.

There are three major reasons:

1. The Problem Size Is Extremely Small :

For 1 process, it was just 80 microseconds.

MPI startup + synchronization overhead alone is often:

- 50–500 microseconds

So what we are measuring is mostly : MPI initialization, MPI_Allreduce latency, Synchronization overhead. But not the actual computation.

When the workload is too small:

$T_{\text{communication}} \gg T_{\text{computation}}$

So adding more processes increases overhead.

2. MPI_Allreduce on Full n Array Every Iteration :

This is expensive:

```
MPI_Allreduce(MPI_IN_PLACE, new_comp.data(), n, MPI_INT, MPI_MIN,
MPI_COMM_WORLD);
```

Cost per iteration: $O(n \log P)$

When n is small \Rightarrow latency dominates

When P increases \Rightarrow communication cost increases

Speedup and Efficiency :

$$S(2) = 0.00008 / 0.000188 \approx 0.42$$

$$S(4) = 0.00008 / 0.0018343 \approx 0.043$$

Parallel efficiency:

$$E(2) = 0.42 / 2 = 0.21 = 21\%$$

$$E(4) = 0.043 / 4 \approx 1\%$$

Estimating maximum graph size processable :

Theoretically, since each iteration scans all edges: T is $I(m+n\log P)$

If time roughly scales linearly with m: T is proportional to m

For 3000 edges, it took about 0.00008 seconds => in 600 seconds, we can do –
 $(600/0.00008)*3000 = \text{22500000000 edges} = 22.5 \text{ B edges}$

2. Communication Profiling:

```
cat: mpi_56436.er: No such file or directory
[cs3401.13@rce problem3_2]$ cat mpi_56436.err
[cs3401.13@rce problem3_2]$ cat mpi_56436.out
=====
SLURM_JOB_ID = 56436
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 0 running on node06.local
Process 1 running on node06.local
Process 2 running on node07.local
Process 3 running on node07.local
Total Execution Time: 0.0020443 seconds
Computation time: 0.000109826
Communication time: 0.00313823
0 0
1 0
2 0
3 0
```

$$\text{Communication-to-Computation Ratio} = 0.00313823 / 0.000109826 = \text{2.85745634}$$

For 1 process :

```
[cs3401.13@rce problem3_2]$ cat mpi_57777.out
=====
SLURM_JOB_ID = 57777
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 0 running on node06.local
Total Execution Time: 8.312e-05 seconds
Computation time: 7.561e-05
Communication time: 5.876e-06
0 0
1 0
2 0
3 0
4 0
```

For 2 processes :

```
=====
SLURM_JOB_ID = 57789
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 0 running on node06.local
Process 1 running on node06.local
Total Execution Time: 0.000165874 seconds
Computation time: 0.000106224
Communication time: 5.6559e-05
0 0
1 0
2 0
3 0
4 0
```

For 4 processes :

```
[cs3401.13@rce problem3_2]$ cat mpi_57790.out
=====
SLURM_JOB_ID = 57790
SLURM_NODELIST = node[06-07]
SLURM_JOB_GPUS =
=====
Process 1 running on node06.local
Process 0 running on node06.local
Process 2 running on node07.local
Process 3 running on node07.local
Total Execution Time: 0.00189802 seconds
Computation time: 0.000103498
Communication time: 0.002706
0 0
1 0
```

As expected in parallel efficiency analysis, the communication time communication cost increases significantly.

From 1 process to 2 processes -> increased by $0.000056559/0.00000587 = 9.635264055x = 10x$ times

From 2 processes to 4 processes -> increased by $0.002706/0.000056559 = 47.843844481x = 48x$ times

=> From 1 process to 4 processes -> increased by 480x times.

If we look at the actual computation time,

From 1 to 2 processes -> $S(2) = 0.00007561/0.000106224 = 0.71179771$

From 2 to 4 processes -> $0.000103498/0.000106224 = 0.97433725$

The main message is MPI_Allreduce(... n integers ...)

Message size per iteration: $n \times \text{sizeof(int)}$

From logs, message size per Allreduce: 0.0038147 MB

3. Computational Complexity:

FLOPs Analysis:

Per edge:

1 min() operation

2 comparison operations

2 assignment operations

~5 operations per edge.

Per iteration: operations = 5m operations

If I iterations: TotalFLOPs = 5mI

Plus: n comparisons for local_changed = nI

So total: $O(I(m+n))$

Communication Complexity:

Each iteration: $O(n \log P)$

So total: $O(I(m+n \log P))$

Theoretical Scaling :

If:

- $m \gg n$
- I constant

Then: $T = O(m)$